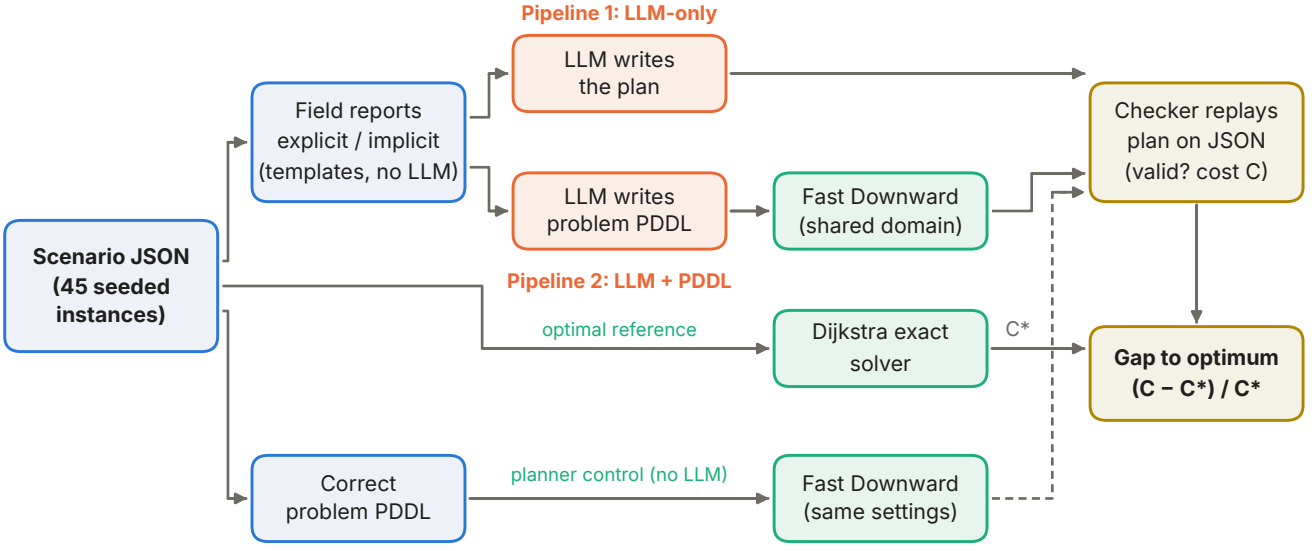
Pipeline 1: LLM-only
LLM writes the plan
Field reports explicit / implicit (templates, no LLM)
LLM writes problem PDDL
Fast Downward (shared domain)
Checker replays plan on JSON (valid? cost C)
Scenario JSON (45 seeded instances)
Pipeline 2: LLM + PDDL
optimal reference
Dijkstra exact solver
C*
Gap to optimum (C − C*) / C*
Correct problem PDDL
planner control (no LLM)
Fast Downward (same settings)